



Deep Learning School

Физтех-Школа Прикладной математики и информатики (ФПМИ) МФТИ

Задание 3

Классификация текстов

В этом задании вам предстоит попробовать несколько методов, используемых в задаче классификации, а также понять насколько хорошо модель понимает смысл слов и какие слова в примере влияют на результат.

```
In [46]: import pandas as pd  
import numpy as np  
import torch
```

```
from torchtext.legacy import datasets

from torchtext.legacy.data import Field, LabelField
from torchtext.legacy.data import BucketIterator

from torchtext.vocab import Vectors, GloVe

import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import random
from tqdm.autonotebook import tqdm

from sklearn.metrics import f1_score
```

```
In [ ]: SEED = 1234

torch.manual_seed(SEED)
torch.backends.cudnn.deterministic = True
```

В этом задании мы будем использовать библиотеку torchtext. Она довольно проста в использовании и поможет нам сконцентрироваться на задаче, а не на написании DataLoader-a.

```
In [87]: TEXT = Field(sequential=True, lower=True, include_lengths=True) # Поле текста
        LABEL = LabelField(dtype=torch.float) # Поле метки
```

Датасет, на котором мы будем проводить эксперименты - это комментарии к фильмам из сайта IMDB.

```
In [88]: train, test = datasets.IMDB.splits(TEXT, LABEL) # загрузим датасет
        train, valid = train.split(random_state=random.seed(SEED)) # разобьем на части
```

```
In [89]: TEXT.build_vocab(train)
        LABEL.build_vocab(train)
```

```
In [90]: device = "cuda" if torch.cuda.is_available() else "cpu"

        train_iter, valid_iter, test_iter = BucketIterator.splits(
            (train, valid, test),
            batch_size = 64,
            sort_within_batch = True,
            device = device)
```

Посмотрим, что выдает итератор

```
In [51]: (text, text_lengths), label = next(iter(train_iter))
```

```
In [52]: text, text.shape
```

```
Out[52]: (tensor([[ 5821,  4376,    2, ...,    9,    9,    2],
                  [191423,  270,  180, ...,  275,  375,   77],
                  [    8,    2,    5, ...,  10,  10,  438],
                  ...,
                  [ 1205,  207,  248, ..., 48564,   29,  197],
                  [   17,   19, 1018, ..., 48815, 59606, 19026],
                  [  487,  672, 33469, ...,    1,    1,    1]]),
          device='cuda:0'), torch.Size([114, 64]))
```

```
In [53]: text_lengths, text_lengths.shape
```

```
Out[53]: (tensor([114, 114, 114, 114, 114, 114, 114, 114, 114, 114, 114, 114, 114, 114,
                  114, 114, 114, 114, 114, 114, 114, 114, 114, 114, 114, 114,
                  114, 114, 114, 113, 113, 113, 113, 113, 113, 113, 113, 113, 113,
                  113, 113, 113, 113, 113, 113, 113, 113, 113, 113, 113, 113, 113,
                  113, 113, 113, 113, 113, 113, 113, 113], device='cuda:0'),
          torch.Size([64]))
```

```
In [54]: label, label.shape
```

```
Out[54]: (tensor([1., 1., 1., 1., 1., 1., 0., 0., 0., 0., 1., 0., 0., 1., 0., 1., 0., 1.,
                  0., 1., 1., 0., 0., 1., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 0., 1.,
                  0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 0., 0.,
                  0., 1., 1., 1., 1., 0., 1., 1., 1., 0.], device='cuda:0'),
          torch.Size([64]))
```

RNN

Для начала попробуем использовать рекуррентные нейронные сети. На семинаре вы познакомились с GRU, вы можете также попробовать LSTM. Можно использовать для классификации как hidden_state, так и output последнего токена.

```
In [55]: class RNNBaseline(nn.Module):
          def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim, n_layers,
                        bidirectional, dropout, pad_idx):

              super().__init__()
              self.num_directions = bidirectional + 1
              self.embedding = nn.Embedding(vocab_size, embedding_dim,
                                             padding_idx = pad_idx).to(device)
              self.rnn = nn.LSTM(embedding_dim, hidden_dim, n_layers, dropout=dropout,
```

```

        bidirectional=bidirectional) # YOUR CODE GOES HERE
self.fc = nn.Linear(hidden_dim * self.num_directions, output_dim) # YOUR CODE GOES HERE
self.dropout = nn.Dropout(dropout)

def forward(self, text, text_lengths):

    #text = [sent len, batch size]

    embedded = self.embedding(text)

    #embedded = [sent len, batch size, emb dim]

    #pack sequence
    packed_embedded = nn.utils.rnn.pack_padded_sequence(embedded, text_lengths.cpu())

    # cell arg for LSTM, remove for GRU
    packed_output, (hidden, cell) = self.rnn(packed_embedded)
    #unpack sequence
    output, output_lengths = nn.utils.rnn.pad_packed_sequence(packed_output)

    #output = [sent len, batch size, hid dim * num directions]
    #output over padding tokens are zero tensors

    #hidden = [num layers * num directions, batch size, hid dim]
    #cell = [num layers * num directions, batch size, hid dim]

    #concat the final forward (hidden[-2,:,:]) and backward (hidden[-1,:,:]) hidden layers
    #and apply dropout

    hidden = torch.cat([hidden[-2], hidden[-1]], dim=1) if self.num_directions == 2 else hidden[-1] # YOUR CODE GOES
    hidden = self.dropout(hidden)

    #hidden = [batch size, hid dim * num directions] or [batch_size, hid dim * num directions]

    return self.fc(hidden).flatten()

```

Поиграйтесь с гиперпараметрами

```

In [74]: vocab_size = len(TEXT.vocab)
emb_dim = 300
hidden_dim = 256
output_dim = 1
n_layers = 4
bidirectional = True

```

```
dropout = 0.2
PAD_IDX = TEXT.vocab.stoi[TEXT.pad_token]
patience=3
# my custom params
max_grad_norm = 0.1
```

```
In [75]: model = RNNBaseline(
        vocab_size=vocab_size,
        embedding_dim=emb_dim,
        hidden_dim=hidden_dim,
        output_dim=output_dim,
        n_layers=n_layers,
        bidirectional=bidirectional,
        dropout=dropout,
        pad_idx=PAD_IDX
    )
```

```
In [76]: model = model.to(device)
```

```
In [77]: opt = torch.optim.Adam(model.parameters())
        loss_func = nn.BCEWithLogitsLoss()

        max_epochs = 20
```

Обучите сетку! Используйте любые вам удобные инструменты, Catalyst, PyTorch Lightning или свои велосипеды.

```
In [78]: import numpy as np

        min_loss = np.inf

        cur_patience = 0

        for epoch in range(1, max_epochs + 1):
            train_loss = 0.0
            model.train()
            pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
            pbar.set_description(f"Epoch {epoch}")
            for it, batch in pbar:
                #YOUR CODE GOES HERE
                opt.zero_grad()
                (text, text_lengths), label = batch
                predict = model(text, text_lengths)
                loss = loss_func(predict, label)
                loss.backward()
```

```

torch.nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
opt.step()
train_loss += loss.item()

train_loss /= len(train_iter)
val_loss = 0.0
model.eval()
pbar = tqdm(enumerate(valid_iter), total=len(valid_iter), leave=False)
pbar.set_description(f"Epoch {epoch}")
for it, batch in pbar:
    # YOUR CODE GOES HERE
    (text, text_lengths), label = batch
    predict = model(text, text_lengths)
    loss = loss_func(predict, label)
    val_loss += loss.item()
val_loss /= len(valid_iter)
if val_loss < min_loss:
    min_loss = val_loss
    best_model = model.state_dict()
else:
    cur_patience += 1
    if cur_patience == patience:
        cur_patience = 0
        break

print('Epoch: {}, Training Loss: {}, Validation Loss: {}'.format(epoch, train_loss, val_loss))
model.load_state_dict(best_model)

```

Epoch: 1, Training Loss: 0.6577718360145597, Validation Loss: 0.6504536700450768

Epoch: 2, Training Loss: 0.5587749980444455, Validation Loss: 0.5076152240320787

Epoch: 3, Training Loss: 0.3531457634411589, Validation Loss: 0.38660544460102664

Epoch: 4, Training Loss: 0.1946592913111196, Validation Loss: 0.4133057494537305

Epoch: 5, Training Loss: 0.08980402345582163, Validation Loss: 0.4782852578466221

Out[78]: <All keys matched successfully>

Посчитайте f1-score вашего классификатора на тестовом датасете.

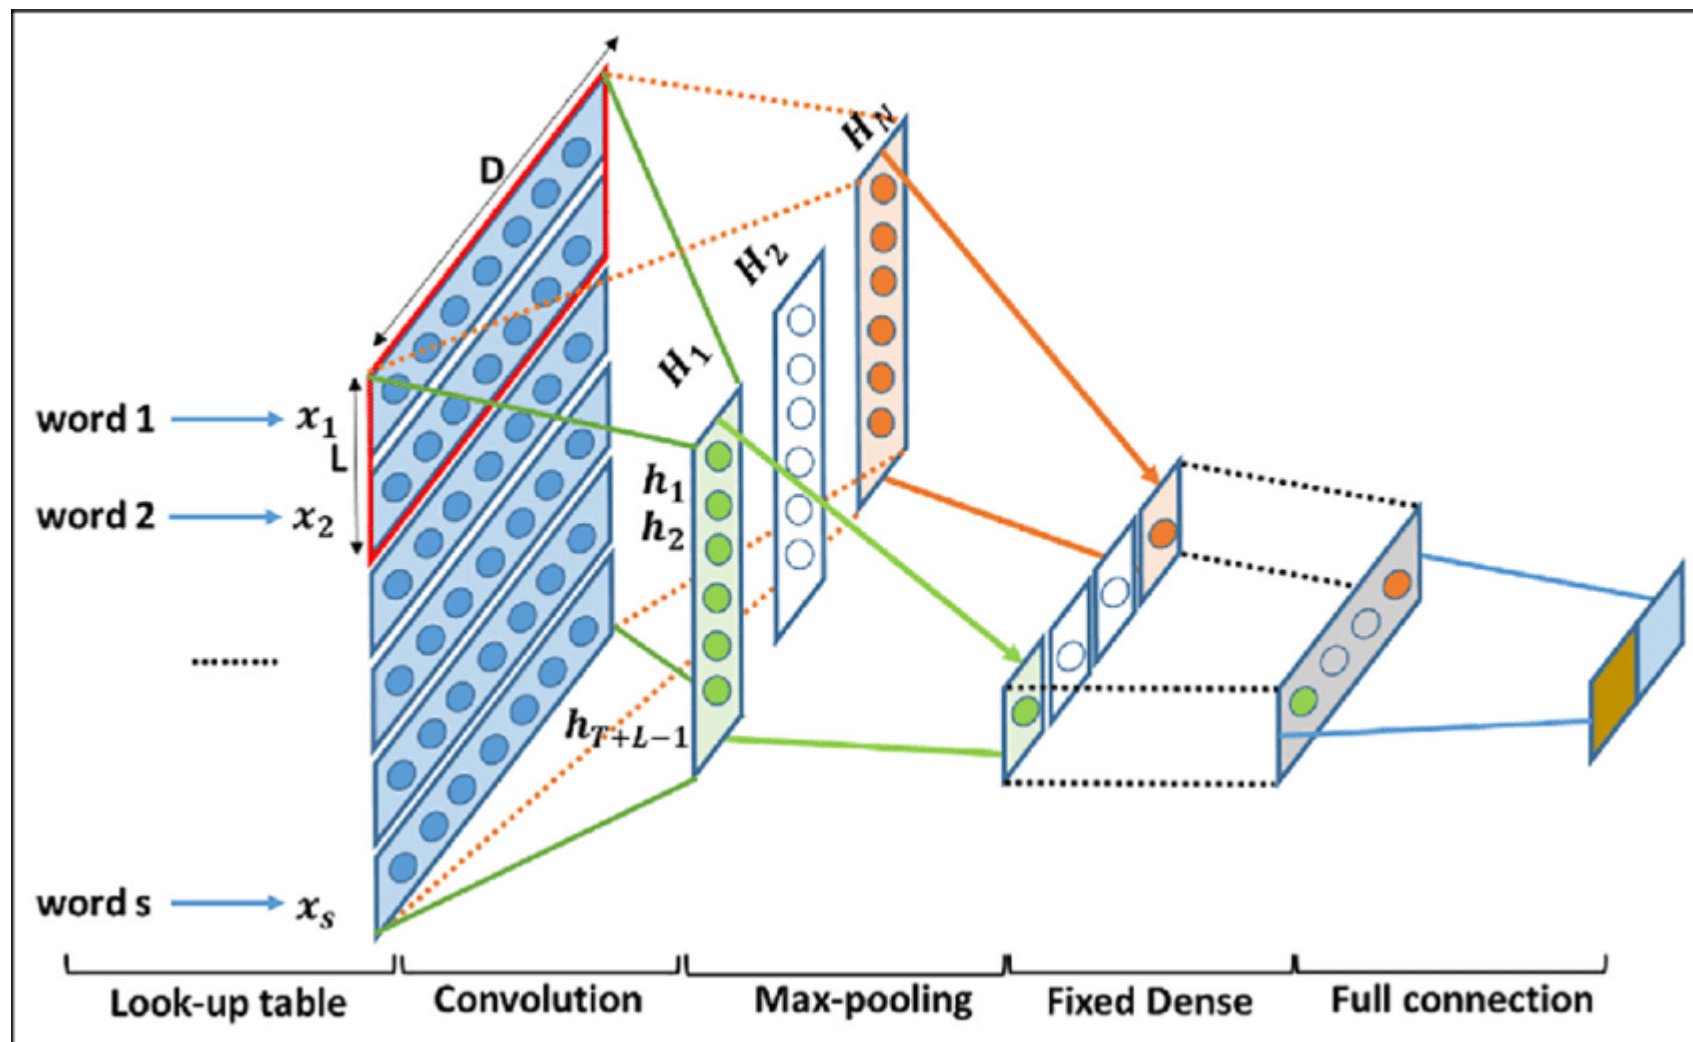
Ответ:

```
In [79]: model.eval()
pbar = tqdm(enumerate(test_iter), total=len(test_iter), leave=False)

predict_labels = []
labels = []
for it, batch in pbar:
    (text, text_lengths), label = batch
    predict_label = model(text, text_lengths) > 0
    predict_labels.extend(predict_label.tolist())
    labels.extend(label.tolist())
f1 = f1_score(labels, predict_labels)
print(f"f1_score = {f1:.4}")
```

f1_score = 0.844

CNN



Для классификации текстов также часто используют сверточные нейронные сети. Идея в том, что как правило сентимент содержат словосочетания из двух-трех слов, например "очень хороший фильм" или "невероятная скука". Проходясь сверткой по этим словам мы получим какой-то большой скор и выхватим его с помощью MaxPool. Далее идет обычная полносвязная сетка. Важный момент: свертки применяются не последовательно, а параллельно. Давайте попробуем!

```
In [100... TEXT = Field(sequential=True, lower=True, batch_first=True) # batch_first тк мы используем conv
LABEL = LabelField(batch_first=True, dtype=torch.float)

train, tst = datasets.IMDB.splits(TEXT, LABEL)
trn, vld = train.split(random_state=random.seed(SEED))
```



```
TEXT.build_vocab(trn)
LABEL.build_vocab(trn)

device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
In [101... train_iter, valid_iter, test_iter = BucketIterator.splits(
    (trn, vld, tst),
    batch_sizes=(128, 256, 256),
    sort=False,
    sort_key= lambda x: len(x.src),
    sort_within_batch=False,
    device=device,
    repeat=False,
)
```

Посмотрим на итератор

```
In [82]: text, label = next(iter(valid_iter))
```

```
In [ ]: text.shape
```

```
Out[ ]: torch.Size([256, 1057])
```

```
In [ ]: label.shape
```

```
Out[ ]: torch.Size([256])
```

Вы можете использовать Conv2d с `in_channels=1, kernel_size=(kernel_sizes[0], emb_dim)` или Conv1d с `in_channels=emb_dim, kernel_size=kernel_size[0]` . Но хорошенько подумайте над shape в обоих случаях.

```
In [103... class CNN(nn.Module):
    def __init__(
        self,
        vocab_size,
        emb_dim,
        out_channels,
        kernel_sizes,
        dropout=0.5,
    ):
        super().__init__()

        self.embedding = nn.Embedding(vocab_size, emb_dim)
```

```

self.conv_0 = nn.Conv1d(emb_dim, out_channels, kernel_size=kernel_sizes[0]) # YOUR CODE GOES HERE
self.conv_1 = nn.Conv1d(emb_dim, out_channels, kernel_size=kernel_sizes[1]) # YOUR CODE GOES HERE
self.conv_2 = nn.Conv1d(emb_dim, out_channels, kernel_size=kernel_sizes[2]) # YOUR CODE GOES HERE
self.fc = nn.Linear(len(kernel_sizes) * out_channels, 1)
self.dropout = nn.Dropout(dropout)

def forward(self, text):

    embedded = self.embedding(text)

    embedded = embedded.permute(0, 2, 1) # may be reshape here

    convded_0 = F.relu(self.conv_0(embedded)) # may be reshape here
    convded_1 = F.relu(self.conv_1(embedded)) # may be reshape here
    convded_2 = F.relu(self.conv_2(embedded)) # may be reshape here

    pooled_0 = F.max_pool1d(convded_0, convded_0.shape[2]).squeeze(2)
    pooled_1 = F.max_pool1d(convded_1, convded_1.shape[2]).squeeze(2)
    pooled_2 = F.max_pool1d(convded_2, convded_2.shape[2]).squeeze(2)

    cat = self.dropout(torch.cat((pooled_0, pooled_1, pooled_2), dim=1))

    return self.fc(cat).flatten()

```

```

In [134... kernel_sizes = [3, 4, 5]
vocab_size = len(TEXT.vocab)
out_channels = 64
dropout = 0.2
dim = 300
patience = 3
# my custom params
max_grad_norm = 1

model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=out_channels,
            kernel_sizes=kernel_sizes, dropout=dropout)

```

```

In [135... model.to(device)

```

```

Out[135... CNN(
  (embedding): Embedding(202268, 300)
  (conv_0): Conv1d(300, 64, kernel_size=(3,), stride=(1,))
  (conv_1): Conv1d(300, 64, kernel_size=(4,), stride=(1,))
  (conv_2): Conv1d(300, 64, kernel_size=(5,), stride=(1,))
  (fc): Linear(in_features=192, out_features=1, bias=True)

```

```
(dropout): Dropout(p=0.2, inplace=False)
)
```

```
In [136... opt = torch.optim.Adam(model.parameters())
loss_func = nn.BCEWithLogitsLoss()
```

```
In [137... max_epochs = 30
```

Обучите!

```
In [138... import numpy as np

min_loss = np.inf

cur_patience = 0

for epoch in range(1, max_epochs + 1):
    train_loss = 0.0
    model.train()
    pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        #YOUR CODE GOES HERE
        opt.zero_grad()
        text, label = batch
        predict = model(text)
        loss = loss_func(predict, label)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
        opt.step()
        train_loss += loss.item()

    train_loss /= len(train_iter)
    val_loss = 0.0
    model.eval()
    pbar = tqdm(enumerate(valid_iter), total=len(valid_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        # YOUR CODE GOES HERE
        text, label = batch
        predict = model(text)
        loss = loss_func(predict, label)
        val_loss += loss.item()
    val_loss /= len(valid_iter)
    if val_loss < min_loss:
```

```
min_loss = val_loss
best_model = model.state_dict()
else:
    cur_patience += 1
    if cur_patience == patience:
        cur_patience = 0
        break

print('Epoch: {}, Training Loss: {}, Validation Loss: {}'.format(epoch, train_loss, val_loss))
model.load_state_dict(best_model)
```

Epoch: 1, Training Loss: 0.571613120119067, Validation Loss: 0.44881325562795005

Epoch: 2, Training Loss: 0.40739470785551696, Validation Loss: 0.3937748700380325

Epoch: 3, Training Loss: 0.3151943307288372, Validation Loss: 0.3639407426118851

Epoch: 4, Training Loss: 0.23513055391555285, Validation Loss: 0.34165724913279216

Epoch: 5, Training Loss: 0.1731998088912372, Validation Loss: 0.33317116151253384

Epoch: 6, Training Loss: 0.12281479493436152, Validation Loss: 0.3312919095158577

Epoch: 7, Training Loss: 0.08288537694589936, Validation Loss: 0.36669970701138177

Epoch: 8, Training Loss: 0.05695953304423903, Validation Loss: 0.35555475354194643

Out[138...] <All keys matched successfully>

Посчитайте f1-score вашего классификатора.

Ответ:

```
In [139...] model.eval()
pbar = tqdm(enumerate(test_iter), total=len(test_iter), leave=False)

predict_labels = []
labels = []
for it, batch in pbar:
    text, label = batch
```



```

input_indices = input_indices.unsqueeze(0)

# input_indices dim: [sequence_length]
seq_length = min_len

# predict
pred = forward_with_sigmoid(input_indices).item()
pred_ind = round(pred)

# generate reference indices for each sample
reference_indices = token_reference.generate_reference(seq_length, device=device).unsqueeze(0)

# compute attributions and approximation delta using layer integrated gradients
attributions_ig, delta = lig.attribute(input_indices, reference_indices, \
                                       n_steps=5000, return_convergence_delta=True)

print('pred: ', LABEL.vocab.itos[pred_ind], '(', '%.2f'%pred, ')', ', delta: ', abs(delta))

add_attributions_to_visualizer(attributions_ig, text, pred, pred_ind, label, delta, vis_data_records_ig)

def add_attributions_to_visualizer(attributions, text, pred, pred_ind, label, delta, vis_data_records):
    attributions = attributions.sum(dim=2).squeeze(0)
    attributions = attributions / torch.norm(attributions)
    attributions = attributions.cpu().detach().numpy()

    # storing couple samples in an array for visualization purposes
    vis_data_records.append(visualization.VisualizationDataRecord(
        attributions,
        pred,
        LABEL.vocab.itos[pred_ind],
        LABEL.vocab.itos[label],
        LABEL.vocab.itos[1],
        attributions.sum(),
        text,
        delta))

```

```

In [142... interpret_sentence(model, 'It was a fantastic performance !', label=1)
interpret_sentence(model, 'Best film ever', label=1)
interpret_sentence(model, 'Such a great show!', label=1)
interpret_sentence(model, 'It was a horrible movie', label=0)
interpret_sentence(model, 'I\'ve never watched something as bad', label=0)
interpret_sentence(model, 'It is a disgusting movie!', label=0)

```

```

pred: pos ( 1.00 ) , delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.01 ) , delta: tensor([0.0001], device='cuda:0', dtype=torch.float64)
pred: pos ( 1.00 ) , delta: tensor([4.7495e-05], device='cuda:0', dtype=torch.float64)

```

```
pred: neg ( 0.02 ) , delta: tensor([9.3134e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.32 ) , delta: tensor([0.0002], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.41 ) , delta: tensor([6.5641e-05], device='cuda:0', dtype=torch.float64)
```

Попробуйте добавить свои примеры!

```
In [143... print('Visualize attributions based on Integrated Gradients')
visualization.visualize_text(vis_data_records_ig)
```

Visualize attributions based on Integrated Gradients

Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (1.00)	pos	1.45	It was a fantastic performance ! pad
pos	neg (0.01)	pos	-0.47	Best film ever pad pad pad pad
pos	pos (1.00)	pos	1.48	Such a great show! pad pad pad
neg	neg (0.02)	pos	-0.06	It was a horrible movie pad pad
neg	neg (0.32)	pos	0.70	I've never watched something as bad pad
neg	neg (0.41)	pos	0.46	It is a disgusting movie! pad pad

Out[143... Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (1.00)	pos	1.45	It was a fantastic performance ! pad
pos	neg (0.01)	pos	-0.47	Best film ever pad pad pad pad
pos	pos (1.00)	pos	1.48	Such a great show! pad pad pad
neg	neg (0.02)	pos	-0.06	It was a horrible movie pad pad
neg	neg (0.32)	pos	0.70	I've never watched something as bad pad
neg	neg (0.41)	pos	0.46	It is a disgusting movie! pad pad

Эмбэдинги слов

Вы ведь не забыли, как мы можем применить знания о word2vec и GloVe. Давайте попробуем!

```
In [151... kernel_sizes = [3, 4, 5]
dropout = 0.2
dim = 300
# my custom params
max_grad_norm = 2

TEXT.build_vocab(trn, vectors=GloVe(dim=dim)) # YOUR CODE GOES HERE
# подсказка: один из импортов пока не использовался, быть может он нужен в строке выше :)
LABEL.build_vocab(trn)

word_embeddings = TEXT.vocab.vectors
vocab_size = len(TEXT.vocab)
```

```
In [152... train, tst = datasets.IMDB.splits(TEXT, LABEL)
trn, vld = train.split(random_state=random.seed(SEED))

device = "cuda" if torch.cuda.is_available() else "cpu"

train_iter, valid_iter, test_iter = BucketIterator.splits(
    (trn, vld, tst),
    batch_sizes=(128, 256, 256),
    sort=False,
    sort_key= lambda x: len(x.src),
    sort_within_batch=False,
    device=device,
    repeat=False,
)
```

```
In [153... model = CNN(vocab_size=vocab_size, emb_dim=dim, out_channels=64,
              kernel_sizes=kernel_sizes, dropout=dropout)

word_embeddings = TEXT.vocab.vectors

prev_shape = model.embedding.weight.shape

model.embedding.weight = nn.Parameter(word_embeddings) # инициализируйте эмбединги

assert prev_shape == model.embedding.weight.shape
model.to(device)

opt = torch.optim.Adam(model.parameters())
```


Вы знаете, что делать.

```
In [154... import numpy as np

min_loss = np.inf

cur_patience = 0

for epoch in range(1, max_epochs + 1):
    train_loss = 0.0
    model.train()
    pbar = tqdm(enumerate(train_iter), total=len(train_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        #YOUR CODE GOES HERE
        opt.zero_grad()
        text, label = batch
        predict = model(text)
        loss = loss_func(predict, label)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
        opt.step()
        train_loss += loss.item()

    train_loss /= len(train_iter)
    val_loss = 0.0
    model.eval()
    pbar = tqdm(enumerate(valid_iter), total=len(valid_iter), leave=False)
    pbar.set_description(f"Epoch {epoch}")
    for it, batch in pbar:
        # YOUR CODE GOES HERE
        text, label = batch
        predict = model(text)
        loss = loss_func(predict, label)
        val_loss += loss.item()
    val_loss /= len(valid_iter)
    if val_loss < min_loss:
        min_loss = val_loss
        best_model = model.state_dict()
    else:
        cur_patience += 1
        if cur_patience == patience:
            cur_patience = 0
            break
```

```
print('Epoch: {}, Training Loss: {}, Validation Loss: {}'.format(epoch, train_loss, val_loss))
model.load_state_dict(best_model)
```

Epoch: 1, Training Loss: 0.44054551672761455, Validation Loss: 0.32145943144957223

Epoch: 2, Training Loss: 0.24896514122068447, Validation Loss: 0.2785116364558538

Epoch: 3, Training Loss: 0.12991056284003885, Validation Loss: 0.27669415970643363

Epoch: 4, Training Loss: 0.049044410922449, Validation Loss: 0.2939944138129552

Epoch: 5, Training Loss: 0.015551697167085253, Validation Loss: 0.3237018610040347

Out[154...] <All keys matched successfully>

Посчитайте f1-score вашего классификатора.

Ответ:

```
In [155...] model.eval()
pbar = tqdm(enumerate(test_iter), total=len(test_iter), leave=False)

predict_labels = []
labels = []
for it, batch in pbar:
    text, label = batch
    predict_label = model(text) > 0
    predict_labels.extend(predict_label.tolist())
    labels.extend(label.tolist())
f1 = f1_score(labels, predict_labels)
print(f"f1_score = {f1:.4}")
```

f1_score = 0.8817

Проверим насколько все хорошо!

```
In [156...] PAD_IND = TEXT.vocab.stoi['pad']

token_reference = TokenReferenceBase(reference_token_idx=PAD_IND)
lig = LayerIntegratedGradients(model, model.embedding)
vis_data_records_ig = []
```

```
interpret_sentence(model, 'It was a fantastic performance !', label=1)
interpret_sentence(model, 'Best film ever', label=1)
interpret_sentence(model, 'Such a great show!', label=1)
interpret_sentence(model, 'It was a horrible movie', label=0)
interpret_sentence(model, 'I\'ve never watched something as bad', label=0)
interpret_sentence(model, 'It is a disgusting movie!', label=0)
```

```
pred: pos ( 0.98 ) , delta: tensor([4.2085e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ) , delta: tensor([5.2274e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.02 ) , delta: tensor([6.0991e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ) , delta: tensor([1.8015e-05], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.25 ) , delta: tensor([0.0002], device='cuda:0', dtype=torch.float64)
pred: neg ( 0.00 ) , delta: tensor([9.9167e-05], device='cuda:0', dtype=torch.float64)
```

```
In [157... print('Visualize attributions based on Integrated Gradients')
visualization.visualize_text(vis_data_records_ig)
```

Visualize attributions based on Integrated Gradients

Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (0.98)	pos	1.87	It was a fantastic performance ! pad
pos	neg (0.00)	pos	0.98	Best film ever pad pad pad pad
pos	neg (0.02)	pos	1.54	Such a great show! pad pad pad
neg	neg (0.00)	pos	0.51	It was a horrible movie pad pad
neg	neg (0.25)	pos	1.90	I've never watched something as bad pad
neg	neg (0.00)	pos	0.49	It is a disgusting movie! pad pad

Out[157... Legend: ☐ Negative ☐ Neutral ☐ Positive

True Label	Predicted Label	Attribution Label	Attribution Score	Word Importance
pos	pos (0.98)	pos	1.87	It was a fantastic performance ! pad
pos	neg (0.00)	pos	0.98	Best film ever pad pad pad pad
pos	neg (0.02)	pos	1.54	Such a great show! pad pad pad
neg	neg (0.00)	pos	0.51	It was a horrible movie pad pad

neg	neg (0.25)	pos	1.90	I've never watched something as bad pad
neg	neg (0.00)	pos	0.49	It is a disgusting movie! pad pad